

```
library(here)
library(pheatmap)
library(corrplot)
```

here() starts at D:/Studies/FINM_34700_MSA/FINM_34700_MSA

corrplot 0.95 loaded

Problem 1

```
path <- here("hw1", "data", "ladyrun25.dat")
ladyrun <- read.table(path)
colnames(ladyrun) <-
c("Country", "100m", "200m", "400m", "800m", "1500m", "3000m", "Marathon")
print(ladyrun)
```

	Country	100m	200m	400m	800m	1500m	3000m	Marathon
1	ARG	11.17	22.94	52.50	2.05	4.20	9.19	150.32
2	AUS	11.11	22.23	48.63	1.97	4.01	8.63	142.60
3	AUT	11.15	22.70	50.62	1.94	4.05	8.78	154.35
4	BEL	11.04	22.20	51.45	1.97	4.08	8.82	143.05
5	BER	11.28	23.05	53.30	2.07	4.29	9.81	174.18
6	BRA	11.17	22.48	50.62	1.97	4.17	9.04	147.41
7	CAN	10.98	22.50	49.91	1.95	3.94	8.54	146.93
8	CHI	11.65	23.84	53.68	2.00	4.22	9.26	152.23
9	CHN	10.79	22.01	45.14	1.93	3.84	8.10	139.65
10	COL	11.31	22.92	49.64	2.04	4.34	9.37	155.19
11	COK	12.52	25.91	61.65	2.28	4.82	11.10	212.33
12	CRC	11.72	23.92	52.57	2.10	4.52	9.84	164.33
13	CZE	11.09	21.97	47.99	1.89	4.03	8.87	145.19
14	DEN	11.42	23.36	52.92	2.02	4.12	8.71	149.34
15	DOM	11.63	23.91	53.02	2.09	4.54	9.89	166.46
16	FIN	11.13	22.39	50.14	2.01	4.10	8.69	148.00
17	FRA	10.73	21.99	48.25	1.94	4.03	8.35	148.27
18	GER	10.81	21.71	47.60	1.92	3.96	8.33	139.32
19	GBR	10.85	22.10	49.43	1.94	3.97	8.37	135.25
20	GRE	10.83	22.67	50.45	2.00	4.09	8.96	153.40
21	GUA	11.92	24.50	55.64	2.15	4.48	9.71	171.33
22	HUN	11.41	23.06	51.50	1.99	4.02	8.55	148.50
23	INA	11.56	23.86	55.08	2.10	4.36	9.50	154.29
24	IND	11.22	22.82	50.79	2.00	4.10	9.11	158.10
25	IRL	11.43	23.02	51.07	2.01	3.98	8.36	142.23
26	ISR	11.27	23.15	52.06	2.07	4.24	9.33	156.36
27	ITA	11.14	22.56	50.30	1.96	3.98	8.59	143.47
28	JPN	11.21	23.33	51.93	2.01	4.16	8.74	139.41
29	KEN	11.47	23.37	51.18	1.97	3.96	8.39	134.07
30	KORS	11.49	23.69	53.67	2.07	4.24	9.01	146.12
31	KORN	11.80	25.10	56.23	1.97	4.25	8.96	145.31

32	LUX	11.76	23.96	56.07	2.07	4.35	9.21	149.23
33	MAS	11.50	23.37	52.56	2.12	4.39	9.31	169.28
34	MRI	11.72	23.83	54.62	2.06	4.33	9.24	167.09
35	MEX	11.09	23.13	48.89	2.02	4.19	8.89	144.06
36	MYA	11.66	23.69	52.96	2.03	4.20	9.08	158.42
37	NED	11.08	22.81	51.35	1.93	4.06	8.31	143.43
38	NZL	11.32	23.13	51.60	1.97	4.10	8.76	146.46
39	NOR	11.10	23.31	52.45	2.03	4.01	8.53	141.06
40	PNG	11.29	23.12	55.18	2.24	4.62	10.21	221.14
41	PHI	11.28	23.01	53.81	2.12	4.41	9.81	165.48
42	POL	10.93	22.13	49.28	1.95	3.99	8.53	144.18
43	POR	11.30	22.88	51.92	1.98	3.96	8.50	143.29
44	ROM	11.30	22.35	49.88	1.92	3.90	8.36	142.50
45	RUS	10.77	21.87	49.11	1.91	3.87	8.38	141.31
46	SAM	12.38	25.45	56.32	2.29	5.42	13.12	191.58
47	SIN	12.13	24.54	55.08	2.12	4.52	9.94	154.41
48	ESP	11.06	22.38	49.67	1.96	4.01	8.48	146.51
49	SWE	11.16	22.82	51.69	1.99	4.09	8.81	150.39
50	SUI	11.34	22.88	51.32	1.98	3.97	8.60	145.51
51	TPE	11.22	22.56	52.74	2.08	4.38	9.63	159.53
52	THA	11.33	23.30	52.60	2.06	4.38	10.07	162.39
53	TUR	11.25	22.71	51.53	2.01	3.92	8.53	144.73
54	USA	10.49	21.34	48.70	1.93	3.92	8.43	139.60

a)

```
print(colMeans(ladyrun[sapply(ladyrun, is.numeric)], na.rm = TRUE))
```

	100m	200m	400m	800m	1500m	3000m
Marathon						
	11.310370	23.070926	51.820185	2.021296	4.186667	9.067222
	153.306852					

Observations

- Scale of 100m, 200m, 400m is in seconds
- Scale of others is in minutes
- Marathon is less interpretable as the distance for marathon is not mentioned

b)

```
# compute covariance matrix
cov_matrix <- cov(ladyrun[sapply(ladyrun, is.numeric)], use =
"complete.obs")
print(cov_matrix)
```

	100m	200m	400m	800m	1500m
3000m					
100m	0.15160363	0.33398644	0.9114414	0.025373096	0.08181258

```


0.23144633
200m      0.33398644 0.82557460  2.2083772 0.060023305 0.19168805
0.53146111
400m      0.91144144 2.20837718  7.4956509 0.191735604 0.55504969
1.56931373
800m      0.02537310 0.06002331  0.1917356 0.007600175 0.02165912
0.06305461
1500m     0.08181258 0.19168805  0.5550497 0.021659119 0.07495849
0.21936604
3000m     0.23144633 0.53146111  1.5693137 0.063054612 0.21936604
0.68175629
Marathon  3.72633326 8.57760297 30.7499534 1.243528686 3.63533836
11.01618166
Marathon
100m      3.726333
200m      8.577603
400m      30.749953
800m      1.243529
1500m     3.635338
3000m     11.016182
Marathon 276.841762

```

```


# compute Pearson correlation matrix
pearson_cor_matrix <- cor(ladyrun[apply(ladyrun, is.numeric)], method
= "pearson", use = "complete.obs")
corrplot(
  pearson_cor_matrix,
  method = "color",          # cleaner filled tiles
  type = "lower",           # lower triangle only
  diag = FALSE,             # hide diagonal
  order = "original",       # keep original column order
  tl.col = "black",         # label color
  tl.srt = 45,              # label angle
  col = colorRampPalette(c("#2166AC", "white", "#B2182B"))(200),
  addCoef.col = "black",    # print correlation values
  number.cex = 0.6,
  mar = c(0, 0, 2, 0),
  title = "Pearson Correlation Heatmap (Lower Triangle)"
)

```




c)

```
# compute Spearman correlation matrix
spearman_cor_matrix <- cor(ladyrun[apply(ladyrun, is.numeric)],
method = "spearman", use = "complete.obs")
corrplot(
  spearman_cor_matrix,
  method = "color",           # cleaner filled tiles
  type = "lower",            # lower triangle only
  diag = FALSE,              # hide diagonal
  order = "original",        # keep original column order
  tl.col = "black",          # label color
  tl.srt = 45,               # label angle
  col = colorRampPalette(c("#2166AC", "white", "#B2182B"))(200),
  addCoef.col = "black",     # print correlation values
  number.cex = 0.6,
  mar = c(0, 0, 2, 0),
  title = "Spearman Correlation Heatmap (Lower Triangle)"
)
```




```
# compute Kendall tau correlation matrix
kendall_cor_matrix <- cor(ladyrun[apply(ladyrun, is.numeric)], method
= "kendall", use = "complete.obs")
corrplot(
  kendall_cor_matrix,
  method = "color",           # cleaner filled tiles
  type = "lower",            # lower triangle only
  diag = FALSE,              # hide diagonal
  order = "original",        # keep original column order
  tl.col = "black",          # label color
  tl.srt = 45,               # label angle
  col = colorRampPalette(c("#2166AC", "white", "#B2182B"))(200),
  addCoef.col = "black",     # print correlation values
  number.cex = 0.6,
  mar = c(0, 0, 2, 0),
  title = "Kendalls tau Correlation Heatmap (Lower Triangle)"
)
```




d)

```
corrplot(
  spearman_cor_matrix - pearson_cor_matrix,
  method = "color",          # cleaner filled tiles
  type = "lower",           # lower triangle only
  diag = FALSE,             # hide diagonal
  order = "original",       # keep original column order
  tl.col = "black",         # label color
  tl.srt = 45,              # label angle
  col = colorRampPalette(c("#2166AC", "white", "#B2182B"))(200),
  addCoef.col = "black",    # print correlation values
  number.cex = 0.6,
  mar = c(0, 0, 2, 0),
  title = "Difference between Spearman and Pearson Correlation"
)
```


Difference between Spearman and Pearson Correlation



```
corrplot(
  kendall_cor_matrix - pearson_cor_matrix,
  method = "color",          # cleaner filled tiles
  type = "lower",           # lower triangle only
  diag = FALSE,             # hide diagonal
  order = "original",       # keep original column order
  tl.col = "black",         # label color
  tl.srt = 45,              # label angle
  col = colorRampPalette(c("#2166AC", "white", "#B2182B"))(200),
  addCoef.col = "black",    # print correlation values
  number.cex = 0.6,
  mar = c(0, 0, 2, 0),
  title = "Difference between Kendall tau and Pearson Correlation"
)
```



```
corrplot(
  kendall_cor_matrix - spearman_cor_matrix,
  method = "color",           # cleaner filled tiles
  type = "lower",            # lower triangle only
  diag = FALSE,              # hide diagonal
  order = "original",        # keep original column order
  tl.col = "black",          # label color
  tl.srt = 45,               # label angle
  col = colorRampPalette(c("#2166AC", "white", "#B2182B"))(200),
  addCoef.col = "black",     # print correlation values
  number.cex = 0.6,
  mar = c(0, 0, 2, 0),
  title = "Difference between Kendall tau and Spearman Correlation"
)
```

Observations

- Kendall's tau has the lowest correlation for all pairs
- Spearman and Pearson are closer together with Spearman being slightly lower in many cases but consistently lower

e)

```
# compute eigenvalues and eigenvectors of pearson correlation matrix
eigen_pearson <- eigen(pearson_cor_matrix)
print(eigen_pearson$values)
print(eigen_pearson$vectors)
```

```
[1] 5.70130407 0.73690632 0.28859891 0.11440672 0.08608808 0.05074016
0.02195575
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[,6]					
[1,]	-0.3720342	-0.4575195	-0.14870245	0.52629124	-0.15450205
[2,]	-0.3738784	-0.4801563	-0.07423786	0.11131548	-0.09164471
[3,]	-0.3747904	-0.3314811	0.48724807	-0.50849863	0.45647911
[4,]	-0.3949123	0.2210770	0.14789147	-0.37710528	-0.76947015

```
[5,] -0.3956582  0.2305757 -0.42485979 -0.13992068  0.08162078 -
0.1431547
[6,] -0.3834289  0.3180749 -0.47659266 -0.07501674  0.37659087
0.1401873
[7,] -0.3490255  0.4970255  0.55267291  0.53351836  0.13707747 -
0.1455350
      [,7]
[1,]  0.08348107
[2,] -0.20389904
[3,]  0.07373480
[4,] -0.15592393
[5,]  0.75036695
[6,] -0.59797109
[7,]  0.03296996
```

f)

```
# sum of eigenvalues should equal the number of variables (trace of
pearson_cor_matrix) (7 in this case)
print(sum(eigen_pearson$values))

[1] 7
```

g)

```
# variance explained by first principal component
variance_explained <- eigen_pearson$values / sum(eigen_pearson$values)
print(variance_explained[1]*100)

[1] 81.4472
```

Observations:

1. The Pearson correlation matrix shows strong positive correlations between the shorter distance events (100m, 200m, 400m) and weaker correlations with the longer distance events (1500m, 3000m, Marathon). This suggests that athletes who perform well in short sprints may not necessarily excel in long-distance events.
2. Correlation of a distance is highest to its closest distances, which is in-line with common expectation.
3. The Spearman and Kendall tau correlation matrices show similar patterns to the Pearson matrix, but with some differences in the strength of correlations. This indicates that while there are monotonic relationships between the events, they may not be perfectly linear.
4. The first principal component explains a significant portion of the variance in the data, indicating that there may be an underlying factor (such as overall athletic ability) that influences performance across all events.

Problem 3

a)

Observations

1) Correlation based PCA is more appropriate as the shorter distances - 100m, 200m, 400m are in seconds and the rest are in minutes. 2) This leads to larger variance for these shorter distances compared to longer distances, whereas intuitively it is common to see larger variances in longer distances.

b)

```
# compute pca of the pearson correlation matrix
pca_pearson <- prcomp(ladyrun[apply(ladyrun, is.numeric)], center =
TRUE, scale. = TRUE)
print(summary(pca_pearson))
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.3877	0.8584	0.53721	0.33824	0.2934	0.22526	0.14817
Proportion of Variance	0.8145	0.1053	0.04123	0.01634	0.0123	0.00725	0.00314
Cumulative Proportion	0.8145	0.9197	0.96097	0.97732	0.9896	0.99686	1.00000


```
# show the factor loadings of PC1 and PC2 to see which variables
contribute most to these components
print(pca_pearson$rotation[, 1:2])
```

	PC1	PC2
100m	0.3720342	0.4575195
200m	0.3738784	0.4801563
400m	0.3747904	0.3314811
800m	0.3949123	-0.2210770
1500m	0.3956582	-0.2305757
3000m	0.3834289	-0.3180749
Marathon	0.3490255	-0.4970255


c)

```
# plot variance explained by each principal component
variance_explained <- pca_pearson$sdev^2 / sum(pca_pearson$sdev^2)
plot(variance_explained, type = "b", xlab = "Principal Component",
ylab = "Proportion of Variance Explained", main = "Variance Explained
by each PCA component")
```

Variance Explained by each PCA component



```
# plot cumulative variance explained by each principal component
plot(cumsum(pca_pearson$sdev^2 / sum(pca_pearson$sdev^2)), type = "b",
     xlab = "Principal Component", ylab = "Cumulative Proportion of
     Variance Explained", main = "Cumulative Variance Explained by PCA")
# add a horizontal line at the variance explained by the first two
# components
abline(h = sum(variance_explained[1:2]), col = "red", lty = 2)
# mark the value on y-axis
text(x = 1, y = sum(variance_explained[1:2]), labels =
     round(sum(variance_explained[1:2]), 2), pos = 1)
```



Observations


1) The number of components to retain depend on the proportion of total variance that we want our analysis to retain. 2) For 90% variance explained, we will choose first two components. 3) For 95% variance explained, we will choose first three components.

d)

```
# plot scores of the observations on the first two principal
components
X_std <- scale(ladyrun[sapply(ladyrun, is.numeric)])

scores = X_std %*% pca_pearson$rotation[, 1:2]

plot(scores, xlab = "PC1", ylab = "PC2", main = "Scores of
Observations on First Two Principal Components")
```



Observations

1) Variance of the scores is higher on PC1 compared to PC2. 2) There is some visible clustering of the scores near origin, with few outliers.

e)

```
pc1 <- pca_pearson$rotation[, 1]
pc2 <- pca_pearson$rotation[, 2]
vars <- rownames(pca_pearson$rotation)
x <- seq_along(vars)

y_all <- c(pc1, pc2)
pad <- 0.1 * diff(range(y_all))

plot(
  x, pc1,
  type = "b", pch = 19, col = "black",
  xaxt = "n",
  ylim = c(min(y_all) - pad, max(y_all) + pad), # increased y-limits
  xlab = "Variables",
  ylab = "Loading",
  main = "PC1 and PC2 Loadings by Variable"
)
```


```

lines(x, pc2, type = "b", pch = 17, col = "red")
axis(1, at = x, labels = vars, las = 2, cex.axis = 0.8)
abline(h = 0, col = "gray60", lty = 2)

# numeric labels near points
text(x, pc1, labels = round(pc1, 2), pos = 3, cex = 0.7, col =
"black")
text(x, pc2, labels = round(pc2, 2), pos = 1, cex = 0.7, col = "red")

legend("topright", legend = c("PC1", "PC2"), col = c("black", "red"),
      lty = 1, pch = c(19, 17), bty = "n")
# ...existing code...

```



Observations:

1) The first PC represents level with almost equal factor loading to all distances. It may be measuring the overall athletic level. 2) The second PC represents slope with positive weight to shorter distances and negative to longer distances. It may be measuring the gradient between shorter and distance times.

f)

```

# compute pca of the covariance matrix
pca_cov <- prcomp(ladyrun[apply(ladyrun, is.numeric)], center = TRUE,

```

```
scale. = FALSE)
print(summary(pca_cov))
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6
PC7						
Standard deviation	16.7670	2.12734	0.54149	0.33417	0.11724	0.05608
0.0273						
Proportion of Variance	0.9827	0.01582	0.00102	0.00039	0.00005	0.00001
0.0000						
Cumulative Proportion	0.9827	0.99852	0.99955	0.99994	0.99999	1.00000
1.0000						


show the factor loadings of PC1 and PC2 to see which variables contribute most to these components

```
print(pca_cov$rotation[, 1:2])
```


	PC1	PC2
100m	0.013596494	0.12014073
200m	0.031347197	0.30270344
400m	0.112057518	0.93290950
800m	0.004483292	0.01248659
1500m	0.013111954	0.03742880
3000m	0.039683998	0.08670969
Marathon	0.992224147	-0.12058732

plot variance explained by each principal component

```
variance_explained <- pca_cov$sdev^2 / sum(pca_cov$sdev^2)
plot(variance_explained, type = "b", xlab = "Principal Component",
ylab = "Proportion of Variance Explained", main = "Variance Explained
by each PCA component")
```

```
# plot cumulative variance explained by each principal component
plot(cumsum(variance_explained), type = "b", xlab = "Principal
Component", ylab = "Cumulative Proportion of Variance Explained", main
= "Cumulative Variance Explained by PCA")
# add a horizontal line at the variance explained by the first two
components
abline(h = sum(variance_explained[1:2]), col = "red", lty = 2)
# mark the value on y-axis
text(x = 1, y = sum(variance_explained[1:2]), labels =
round(sum(variance_explained[1:2]), 2), pos = 1)
```




Observations

1) Almost all of the variance is explained by the first PC

```
# plot scores of the observations on the first two principal components
X_demean <- scale(ladyrun[sapply(ladyrun, is.numeric)], scale=FALSE)

scores = X_demean %*% pca_cov$rotation[, 1:2]

plot(scores, xlab = "PC1", ylab = "PC2", main = "Scores of
Observations on First Two Principal Components")
```



Observations:

1) The scale of PC1 is much larger than PC2, which indicates that PC1 explains much more variance than PC2. This is consistent with the variance explained plot where the first component explains a large portion of the variance.

```
pc1 <- pca_cov$rotation[, 1]
pc2 <- pca_cov$rotation[, 2]
vars <- rownames(pca_cov$rotation)
x <- seq_along(vars)

y_all <- c(pc1, pc2)
pad <- 0.1 * diff(range(y_all))

plot(
  x, pc1,
  type = "b", pch = 19, col = "black",
  xaxt = "n",
  ylim = c(min(y_all) - pad, max(y_all) + pad), # increased y-limits
  xlab = "Variables",
  ylab = "Loading",
  main = "PC1 and PC2 Loadings by Variable"
)
```


```

lines(x, pc2, type = "b", pch = 17, col = "red")
axis(1, at = x, labels = vars, las = 2, cex.axis = 0.8)
abline(h = 0, col = "gray60", lty = 2)

# numeric labels near points
text(x, pc1, labels = round(pc1, 2), pos = 3, cex = 0.7, col =
"black")
text(x, pc2, labels = round(pc2, 2), pos = 1, cex = 0.7, col = "red")

legend("topright", legend = c("PC1", "PC2"), col = c("black", "red"),
      lty = 1, pch = c(19, 17), bty = "n")

```



Observations:

1) PC1 has maximum factor loadings on 400m and second largest on 200m and negative loadings on Marathon. 2) PC2 has maximum factor loadings on Marathon.

g)

Observations:

1) PCA conclusions are highly sensitive to the choice of covariance vs correlation matrix (standardization) when the scale of all the variables is not the same. 2) PCA on Correlation matrix is the correct approach in our case.

Problem 4)

```
path <- here("hw1", "data", "treasury_yield_panel.csv")
treasury_yield_panel <- read.csv(path)
```

```
print(head(treasury_yield_panel))
```

	date	X1Y	X2Y	X3Y	X5Y	X7Y	X10Y	X20Y	X30Y
1	2014-01-02	0.13	0.39	0.76	1.72	2.41	3.00	3.68	3.92
2	2014-01-03	0.13	0.41	0.80	1.73	2.42	3.01	3.69	3.93
3	2014-01-06	0.12	0.40	0.78	1.70	2.38	2.98	3.66	3.90
4	2014-01-07	0.13	0.40	0.80	1.69	2.37	2.96	3.64	3.88
5	2014-01-08	0.13	0.43	0.87	1.77	2.44	3.01	3.67	3.90
6	2014-01-09	0.13	0.44	0.86	1.75	2.41	2.97	3.62	3.88

a)

```
# Calculate daily changes in yields in bps for each maturity
```

```
yield_changes_bps <- treasury_yield_panel
```

```
num_cols <- sapply(yield_changes_bps, is.numeric)
```

```
k <- 1
```

```
yield_changes_bps[num_cols] <- lapply(
  yield_changes_bps[num_cols],
  function(x) c(rep(NA, k), diff(x, lag = k))
)
```

```
yield_changes_bps[num_cols] <- 1e4 * yield_changes_bps[num_cols]
```

```
yield_changes_bps <-
```

```
yield_changes_bps[complete.cases(yield_changes_bps), ]
```

```
print(head(yield_changes_bps))
```

	date	X1Y	X2Y	X3Y	X5Y	X7Y	X10Y	X20Y	X30Y
2	2014-01-03	0	200	400	100	100	100	100	100
3	2014-01-06	-100	-100	-200	-300	-400	-300	-300	-300
4	2014-01-07	100	0	200	-100	-100	-200	-200	-200
5	2014-01-08	0	300	700	800	700	500	300	200
6	2014-01-09	0	100	-100	-200	-300	-400	-500	-200
7	2014-01-10	-100	-500	-900	-1100	-1200	-900	-800	-800

b)

```
# apply pca to the centered yield changes in bps
```

```
pca_yield_changes <-
```

```
prcomp(yield_changes_bps[sapply(yield_changes_bps, is.numeric)],
  center = TRUE, scale. = FALSE)
```

```
print(summary(pca_yield_changes))
```

```

Importance of components:
      PC1      PC2      PC3      PC4      PC5
PC6
Standard deviation 1319.1917 476.5249 213.60352 132.53831 94.26186
68.36535
Proportion of Variance 0.8482 0.1107 0.02224 0.00856 0.00433
0.00228
Cumulative Proportion 0.8482 0.9589 0.98112 0.98968 0.99401
0.99629
      PC7      PC8
Standard deviation 66.95945 55.93084
Proportion of Variance 0.00219 0.00152
Cumulative Proportion 0.99848 1.00000

print(pca_yield_changes$rotation[, 1:3])
      PC1      PC2      PC3
X1Y 0.2056368 -0.43464427 0.789147203
X2Y 0.3321860 -0.46305493 0.002025498
X3Y 0.3723304 -0.33858862 -0.250518251
X5Y 0.4033192 -0.12038526 -0.329321447
X7Y 0.4098169 0.06141526 -0.235730411
X10Y 0.3850044 0.22025556 -0.052376753
X20Y 0.3480298 0.42840217 0.215198015
X30Y 0.3297977 0.48136977 0.318446610

# variance explained by first three principal components
variance_explained <- pca_yield_changes$sdev^2 /
sum(pca_yield_changes$sdev^2)
print(variance_explained[1:3]*100)

[1] 84.82045 11.06766 2.22383

```

c)

```

xc = scale(yield_changes_bps[apply(yield_changes_bps, is.numeric)],
center = TRUE, scale = FALSE)

v1 = pca_yield_changes$rotation[, 1]
v2 = pca_yield_changes$rotation[, 2]

z1 = xc %*% v1
z2 = xc %*% v2

xc_2 = z1 %*% t(v1) + z2 %*% t(v2)

```

c)(i)

```

r = xc - xc_2

```

```

print(head(r %*% v1))
print(tail(r %*% v1))

      [,1]
2  1.598721e-14
3  1.207923e-13
4  3.108624e-14
5 -2.273737e-13
6  6.394885e-14
7  2.629008e-13
      [,1]
3040 -4.014566e-13
3041 -1.598721e-14
3042 -9.570122e-14
3043 -4.973799e-14
3044 -3.552714e-15
3045 -2.486900e-14

print(sum((r %*% v1) > 1e-10))

[1] 0

print(head(r %*% v2))
print(tail(r %*% v2))

      [,1]
2  1.438849e-13
3  8.881784e-14
4  2.629008e-13
5  1.563194e-13
6  3.979039e-13
7  2.486900e-13
      [,1]
3040  4.654055e-13
3041  2.238210e-13
3042  2.648992e-13
3043  9.947598e-14
3044 -6.181722e-13
3045  4.121148e-13

print(sum((r %*% v2) > 1e-10))

[1] 0

```

c)(ii)

```

# find norm of each residual row
residual_norms <- apply(r, 1, function(row) sqrt(sum(row^2)))
print(head(residual_norms))
print(tail(residual_norms))

```

	2	3	4	5	6	7
	271.2101	122.2386	201.3597	556.6132	271.2502	468.9174

	3040	3041	3042	3043	3044	3045
	134.7054	243.2861	144.5018	224.9579	252.9636	129.2823

```
# find index of the max of the residual norms
max_residual_norm <- max(residual_norms)
max_residual_index <- which(residual_norms == max_residual_norm)
print(max_residual_index)

2299
2298

# plot xc, xc_2, and r for the observation with the largest residual
norm for numeric columns
index_to_plot <- max_residual_index[1] # in case there are multiple,
just take the first one
date <- yield_changes_bps[index_to_plot, 'date']

vars <- colnames(xc)
x <- seq_along(vars)


y_all <- c(xc[index_to_plot, ], xc_2[index_to_plot, ])
pad <- 0.1 * diff(range(y_all))

plot(
  x, xc[index_to_plot, ],
  type = "b", pch = 19, col = "black",
  xaxt = "n",
  ylim = c(min(y_all) - pad, max(y_all) + pad), # increased y-limits
  xlab = "Variables",
  ylab = "Loading",
  main = paste("Plotting for date:", date)
)

lines(x, xc_2[index_to_plot, ], type = "b", pch = 17, col = "red")
axis(1, at = x, labels = vars, las = 2, cex.axis = 0.8)
abline(h = 0, col = "gray60", lty = 2)

# numeric labels near points
text(x, xc[index_to_plot, ], labels = round(xc[index_to_plot, ], 2),
pos = 3, cex = 0.7, col = "black")
text(x, xc_2[index_to_plot, ], labels = round(xc_2[index_to_plot, ],
2), pos = 1, cex = 0.7, col = "red")

legend("topright", legend = c("original", "reconstructed from
PC1,PC2"), col = c("black", "red"), lty = 1, pch = c(19, 17), bty =
"n")
```

Observations:

1) The level moves are largely captured by the two component reconstruction. 2) The slope is also largely captured by the two component reconstruction. 3) The curvature is not captured by the first two components as we can visibly note that the lower and higher terms are reduced less in comparison to the original and medium terms are reduced more in comparison to the original.

Problem 5)

a)

```
# fit a two factor factor-analysis model to ladyrun using factanal
data <- ladyrun[sapply(ladyrun, is.numeric)]
factanal_result <- factanal(data, factors = 2)
print(factanal_result)
```

Call:

```
factanal(x = data, factors = 2)
```

Uniquenesses:

100m	200m	400m	800m	1500m	3000m	Marathon
0.097	0.009	0.194	0.155	0.018	0.037	0.324

Loadings:

	Factor1	Factor2
100m	0.411	0.856
200m	0.375	0.922
400m	0.425	0.790
800m	0.762	0.514
1500m	0.865	0.483
3000m	0.894	0.405
Marathon	0.761	0.311

	Factor1	Factor2
SS loadings	3.198	2.966
Proportion Var	0.457	0.424
Cumulative Var	0.457	0.881

Test of the hypothesis that 2 factors are sufficient.
The chi square statistic is 33.71 on 8 degrees of freedom.
The p-value is 4.58e-05

plot the factor loadings of the two factors

```
loadings <- factanal_result$loadings[, 1:2]
```

```
vars <- colnames(data)
```

```
x <- seq_along(vars)
```

```
y_all <- c(loadings[, 1], loadings[, 2])
```

```
pad <- 0.1 * diff(range(y_all))
```

```
plot(
  x, loadings[, 1],
  type = "b", pch = 19, col = "black",
  xaxt = "n",
  ylim = c(min(y_all) - pad, max(y_all) + pad), # increased y-limits
  xlab = "Variables",
  ylab = "Loading",
  main = "Factor Loadings for Factor 1 and Factor 2"
)
```

```
lines(x, loadings[, 2], type = "b", pch = 17, col = "red")
```

```
axis(1, at = x, labels = vars, las = 2, cex.axis = 0.8)
```


```
abline(h = 0, col = "gray60", lty = 2)
```

numeric labels near points

```
text(x, loadings[, 1], labels = round(loadings[, 1], 2), pos = 3, cex = 0.7, col = "black")
```

```
text(x, loadings[, 2], labels = round(loadings[, 2], 2), pos = 1, cex = 0.7, col = "red")
```

```
legend("topright", legend = c("Factor 1", "Factor 2"), col = c("black", "red"), lty = 1, pch = c(19, 17), bty = "n")
```



b)

Observations:

1) Factor 1 loads the longer distances more strongly. 2) Factor 2 loads the shorter distances more strongly. 3) Marathon seems to have the largest uniqueness, followed by 400m and 800m (in the order).

c)

Observations:

1) Similarity: Two factor Factor-analysis (trying to capture common variance) explained about 90% of the variance similar to two factor PCA (trying to capture total variance). 2) Difference: The first PC component captured the maximum variance (90%) in the data, while the first factor captured less than 50% of the variance in the data.
